

Encode latency vs model prefill - the codec is not the bottleneck

